

```
Activities Terminal Sat 03:08
hduser@ubuntu: /usr/local/apache-hive-2.1.0-bin/bin

File Edit View Search Terminal Help
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Logging initialized using configuration in jar:file:/usr/local/apache-hive-2.1.0-bin/lib/hive-common-2.1.0.jar!/hive-log4j2.
properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution eng
ine (i.e. spark, tez) or using Hive 1.X releases.
hive> CREATE DATABASE IF NOT EXISTS lab9 COMMENT 'employee program' WITH DBPROPERTIES ('creator'='CHANDANA');
OK
Time taken: 3.886 seconds
hive> SHOW DATABASES;
OK
default
lab9
Time taken: 1.768 seconds, Fetched: 2 row(s)
hive> DESCRIBE DATABASE lab9;
OK
lab9 hive employee program hdfs://localhost:54310/user/hive/warehouse/lab9.db hduser USER
Time taken: 0.104 seconds, Fetched: 1 row(s)
hive> USE lab9;
OK
Time taken: 0.08 seconds
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Employee(EmpID INT,EmpName STRING,Designation STRING,Salary FLOAT) ROW FORMAT DELI
MITED FIELDS TERMINATED BY '\t';
OK
Time taken: 1.98 seconds
hive> LOAD DATA LOCAL INPATH '/home/chandana/Desktop/employeeInput.txt' OVERWRITE INTO TABLE Employee;
Loading data to table lab9.employee
OK
Time taken: 5.339 seconds
hive>
```

```
Activities Terminal Sat 03:13
hduser@ubuntu: /usr/local/apache-hive-2.1.0-bin/bin

File Edit View Search Terminal Help
OK
Time taken: 3.119 seconds
hive> SELECT * FROM Employee;
OK
1 John Manager 50000.0
2 Kelly Sr. Manager 60000.0
3 Harry Software Engineer 45000.0
4 Kedrik Software Engineer 30000.0
5 Tom Test Engineer 20000.0
6 Ubern Test Engineer 20000.0
7 Rohan HR manager 25000.0
8 Joana Software Engineer 23000.0
Time taken: 2.62 seconds, Fetched: 8 row(s)
hive> CREATE VIEW emp_30000 AS SELECT * FROM Employee WHERE Salary>30000;
OK
Time taken: 0.386 seconds
hive> SELECT * FROM emp_30000;
OK
1 John Manager 50000.0
2 Kelly Sr. Manager 60000.0
3 Harry Software Engineer 45000.0
Time taken: 0.825 seconds, Fetched: 3 row(s)
hive> ALTER TABLE Employee ADD COLUMNS(DeptID INT);
OK
Time taken: 0.391 seconds
hive> LOAD DATA LOCAL INPATH '/home/chandana/Desktop/employeeInputAltered.txt' OVERWRITE INTO TABLE Employee;
Loading data to table default.employee
OK
Time taken: 0.631 seconds
hive>
```

```
Activities Terminal Sat 03:15
hduser@ubuntu: /usr/local/apache-hive-2.1.0-bin/bin

File Edit View Search Terminal Help

hive> SELECT * FROM Employee ORDER BY DeptID;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hduser_20201219031430_f37e5e18-d55b-4d50-8836-b3d86a9ca242
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (Local Hadoop)
2020-12-19 03:14:36,268 Stage-1 map = 0%, reduce = 0%
2020-12-19 03:14:45,727 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local426704584_0001
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 1944 HDFS Write: 972 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
2      Kelly      Sr. Manager      60000.0 11
1      John       Manager 50000.0 11
7      Rohan      HR manager       25000.0 12
8      Joana      Software Engineer 23000.0 13
4      Kedrik     Software Engineer 30000.0 13
3      Harry      Software Engineer 45000.0 13
6      Ubern      Test Engineer    20000.0 15
5      Tom        Test Engineer    20000.0 15
Time taken: 15.476 seconds, Fetched: 8 row(s)
```

```
hive> SELECT DeptID,count(*) FROM Employee WHERE Salary>=30000 GROUP BY DeptID;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hduser_20201219031553_dd2c795c-2640-4463-8c16-298f28a793a1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (Local Hadoop)
2020-12-19 03:15:56,064 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local518152724_0002
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 2454 HDFS Write: 972 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
11      2
13      2
Time taken: 2.915 seconds, Fetched: 2 row(s)
```

```
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Job running in-process (Local Hadoop)
2020-12-19 03:18:13,778 Stage-3 map = 100%, reduce = 0%
Ended Job = job_local1327814845_0003
MapReduce Jobs Launched:
Stage-Stage-3: HDFS Read: 1542 HDFS Write: 546 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
1      John       Manager 50000.0 Business Management
2      Kelly      Sr. Manager      60000.0 Business Management
3      Harry      Software Engineer 45000.0 Development
4      Kedrik     Software Engineer 30000.0 Development
5      Tom        Test Engineer    20000.0 Testing
6      Ubern      Test Engineer    20000.0 Testing
7      Rohan      HR manager       25000.0 HR
8      Joana      Software Engineer 23000.0 Development
Time taken: 51.043 seconds, Fetched: 8 row(s)
```